

Deploy Kafka, Zookeeper, kafka Connect and Spark containers in GCP using Kubernetes

Initial Setup:	2
Prepare and push Images:.....	2
Create a cluster in Google Cloud:	2
Deployment Templates:.....	3
Deploying service:.....	5
Testing our deployment:.....	6

Initial Setup:

Visit kubernetes engine page and select the project (This will take several mins to start kubernetes engine)

https://console.cloud.google.com/projectselector/kubernetes?_ga=2.201928657.-1707404544.1516051830

Once kubernetes engine is ready:

```
gcloud components install kubectl
```

Prepare and push Images:

Check docker images that are locally stored and tag them to be pushed to google cloud:. We are interested in midsw/spark/kafka/zookeeper

[kurapati-new:~] kurapati% docker images

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
midsw205/base	0.1.8	ee33b8fcc42c	4 weeks ago	2.68GB
midsw205/spark-python	0.0.5	0756ef14312d	5 weeks ago	3.17GB
confluentinc/cp-kafka	latest	8fa6da41c4ae	6 weeks ago	535MB
confluentinc/cp-zookeeper	latest	2acdb712eee3	6 weeks ago	535MB

NOTE : I trimmed the above output. You may have more entries, but we are interested in four images for now. Let's now tag them and push

```
[kurapati-new:~] kurapati% docker tag ee33b8fcc42c gcr.io/w205-1/mids
[kurapati-new:~] kurapati% docker tag 0756ef14312d gcr.io/w205-1/spark-python
[kurapati-new:~] kurapati% docker tag 8fa6da41c4ae gcr.io/w205-1/kafka
[kurapati-new:~] kurapati% docker tag 2acdb712eee3 gcr.io/w205-1/zookeeper
[kurapati-new:~] kurapati%
[kurapati-new:~] kurapati% gcloud docker -- push gcr.io/w205-1/mids
[kurapati-new:~] kurapati% gcloud docker -- push gcr.io/w205-1/spark-python
[kurapati-new:~] kurapati% gcloud docker -- push gcr.io/w205-1/kafka
[kurapati-new:~] kurapati% gcloud docker -- push gcr.io/w205-1/zookeeper
```

Create a cluster in Google Cloud:

I am creating a cluster with 5 nodes below (5 nodes is the maximum in n1-standard-1 flavor). You can change the zone based on where you live.

```
[kurapati-new:~] kurapati% gcloud container clusters create kafka --num-nodes=5 --zone northamerica-northeast1-a
```

Check if the cluster is created and the computes are operational:

```
[kurapati-new:~] kurapati% gcloud compute instances list
```

NAME	ZONE	MACHINE_TYPE	PREEMPTIBLE	INTERNAL_IP	EXTERNAL_IP	STATUS
gke-kafka-default-pool-bad8e9fd-2pnb	northamerica-northeast1-a	n1-standard-1		10.162.0.2	35.203.13.30	RUNNING
gke-kafka-default-pool-bad8e9fd-9jc6	northamerica-northeast1-a	n1-standard-1		10.162.0.6	35.203.7.100	RUNNING
gke-kafka-default-pool-bad8e9fd-mrk1	northamerica-northeast1-a	n1-standard-1		10.162.0.3	35.203.10.111	RUNNING
gke-kafka-default-pool-bad8e9fd-mtrm	northamerica-northeast1-a	n1-standard-1		10.162.0.4	35.203.3.234	RUNNING
gke-kafka-default-pool-bad8e9fd-w4t1	northamerica-northeast1-a	n1-standard-1		10.162.0.5	35.203.10.53	RUNNING

```
[kurapati-new:~/W205/flask-with-kafka-and-spark] kurapati%
```

Deployment Templates:

Now, let us get to creating deployment templates to spin up our containers. Since we have docker-compose already we need to convert them into kubernetes deployment template (Yes, they both are different!). Easiest way is to use a fantastic tool called compose. Refer to <http://kompose.io/> for more details of the project.

You need to grab the Kompose binary from the above link.

Snippet for MAC :

```
curl -L https://github.com/kubernetes/kompose/releases/download/v1.11.0/kompose-darwin-amd64 -o kompose
```

```
chmod +x kompose  
sudo mv ./kompose /usr/local/bin/kompose
```

Once downloaded, run the tool from the directory where docker-compose.yaml is present :

```
kompose convert
```

This should generate a bunch of service and deployment templates.

Here are the files on my computer:

```
[kurapati-new:~/W205/flask-with-kafka-and-spark] kurapati% ls *deployment*
```

kafka-deployment.yaml mids-deployment.yaml spark-
deployment.yaml zookeeper-deployment.yaml

```
[kurapati-new:~/W205/flask-with-kafka-and-spark] kurapati% ls *service*  
kafka-service.yaml   mids-service.yaml   spark-service.yaml   zookeeper-service.yaml
```

```
[kurapati-new:~/W205/flask-with-kafka-and-spark] kurapati% ls *volume*  
mids-claim0-persistentvolumeclaim.yaml   spark-claim0-persistentvolumeclaim.yaml  
[kurapati-new:~/W205/flask-with-kafka-and-spark] kurapati%
```

For MIDS and spark containers, I had local volume mounts. In GCP, this can be ignored for the moment but you can create the volumes as well if you desire to mount files from host to container.

Sample deployment and service templates for zookeeper is below:

```
[kurapati-new:~/W205/flask-with-kafka-and-spark] kurapati% more zookeeper-  
deployment.yaml
```

```
apiVersion: extensions/v1beta1  
kind: Deployment  
metadata:  
  annotations:  
    kompose.cmd: /usr/local/bin/kompose convert  
    kompose.version: 1.11.0 (39ad614)  
  creationTimestamp: null  
  labels:  
    io.kompose.service: zookeeper  
  name: zookeeper  
spec:  
  replicas: 1  
  strategy: {}  
  template:  
    metadata:  
      creationTimestamp: null  
      labels:  
        io.kompose.service: zookeeper  
    spec:  
      containers:  
      - env:  
        - name: ZOOKEEPER_CLIENT_PORT  
          value: "32181"  
        - name: ZOOKEEPER_TICK_TIME  
          value: "2000"  
        image: confluentinc/cp-zookeeper:latest
```

```
  name: zookeeper
  resources: {}
  restartPolicy: Always
status: {}
```

[kurapati-new:~/W205/flask-with-kafka-and-spark] kurapati% more zookeeper-service.yaml

```
apiVersion: v1
kind: Service
metadata:
  annotations:
    kompose.cmd: /usr/local/bin/kompose convert
    kompose.version: 1.11.0 (39ad614)
  creationTimestamp: null
  labels:
    io.kompose.service: zookeeper
  name: zookeeper
spec:
  clusterIP: None
  ports:
  - name: headless
    port: 55555
    targetPort: 0
  selector:
    io.kompose.service: zookeeper
status:
  loadBalancer: {}
```

Deploying service:

Now let us bring up the service and deployments in using kubectl using below commands

```
kubectl create --filename zookeeper-deployment.yaml
kubectl create --filename zookeeper-service.yaml
kubectl create --filename kafka-deployment.yaml
kubectl create --filename kafka-service.yaml
kubectl create --filename spark-service.yaml
kubectl create --filename spark-deployment.yaml
kubectl create --filename spark-claim0-persistentvolumeclaim.yaml [optional]
kubectl create --filename mids-service.yaml
kubectl create --filename mids-deployment.yaml
kubectl create --filename mids-claim0-persistentvolumeclaim.yaml [optional]
```

```
[kurapati-new:~] kurapati% kubectl get pods -o wide
```

Testing our deployment:

Processed a total of 42 messages

Login to Spark container and fetch messages through pyspark.

```
[kurapati-new:~/W205/flask-with-kafka-and-spark] kurapati% kubectl exec -it spark-654964b959-r8sqp bash
```

```
root@spark-654964b959-r8sqp:/spark-2.2.0-bin-hadoop2.6#
```

```
root@spark-654964b959-r8sqp:/spark-2.2.0-bin-hadoop2.6# pyspark
Python 3.6.1 |Anaconda 4.4.0 (64-bit)| (default, May 11 2017, 13:09:58)
[GCC 4.4.7 20120313 (Red Hat 4.4.7-1)] on linux
```

```
>>> numbers = spark \
...   .read \
...   .format("kafka") \
...   .option("kafka.bootstrap.servers", "kafka:29092") \
...   .option("subscribe", "events") \
...   .option("startingOffsets", "earliest") \
...   .option("endingOffsets", "latest") \
...   .load()
>>> numbers_as_strings=numbers.selectExpr("CAST(key AS STRING)", "CAST(value AS STRING)")
>>> numbers_as_strings.show()
+----+-----+
| key|value|
+----+-----+
|null|  1|
|null|  2|
|null|  3|
..
|null| 19|
|null| 20|
+----+-----+
only showing top 20 rows

>>>
```

Enjoy the cloud!

References:

<https://cloud.google.com/kubernetes-engine/docs/tutorials/hello-app>

<https://kubernetes.io/docs/admin/cluster-large/>